

Fetal Health Classification

Import libraries to be used

```
In [ ]:

import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
from matplotlib import cm
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import StandardScaler
from sklearn.feature_selection import SelectKBest
from sklearn.feature_selection import f_classif
from sklearn.model_selection import cross_val_score
from sklearn.metrics import classification_report, confusion_matrix
from sklearn.model_selection import GridSearchCV, StratifiedKFold
```

Exploratory Data Analysis and Pre-Processing

```
In [ ]:

data = pd.read_csv('fetal_health.csv')
data.info()
data.head()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 2126 entries, 0 to 2125
Data columns (total 22 columns):
#   Column
---  ---
0   baseline value
1   accelerations
2   fetal_movement
3   uterine_contractions
4   light_decelerations
5   severe_decelerations
6   prolonged_decelerations
7   abnormal_short_term_variability
8   mean_value_of_short_term_variability
9   percentage_of_time_with_abnormal_long_term_variability
10  mean_value_of_long_term_variability
11  histogram_width
12  histogram_min
13  histogram_max
14  histogram_number_of_peaks
15  histogram_number_of_zeroes
16  histogram_mode
17  histogram_mean
18  histogram_median
19  histogram_variance
20  histogram_tendency
21  fetal_health
dtypes: float64(22)
memory usage: 365.5 KB
```

Out[]:

	baseline value	accelerations	fetal_movement	uterine_contractions	light_decelerations	severe_decelerations	prolongued_deceler
0	120.0	0.000	0.0	0.000	0.000	0.0	
1	120.0	0.000	0.0	0.000	0.000	0.0	

1	132.0	0.000	0.0	0.000	0.000	0.0	
2	baseline value	accelerations	fetal_movement	uterine_contractions	light_decelerations	severe_decelerations	prolongued_decelerations
3	134.0	0.003	0.0	0.008	0.003	0.0	
4	132.0	0.007	0.0	0.008	0.000	0.0	

Note that there are no null data values and all column Dtypes are floats. There are 21 columns which translates to 21 features to analyze. Let's now remove duplicate data entries.

In []:

```
unduplicated_data = data.copy()
unduplicated_data.drop_duplicates(inplace=True)
print("Removed " + str(data.shape[0] - unduplicated_data.shape[0]) + " duplicates")
data = unduplicated_data
```

Removed 13 duplicates

In []:

```
print("Total: " + str(data.shape[0]) + " samples")
```

Total: 2113 samples

In []:

```
data.describe().T
```

Out[]:

	count	mean	std	min	25%	50%	75%	max
baseline value	2113.0	133.304780	9.837451	106.0	126.000	133.000	140.000	160.000
accelerations	2113.0	0.003188	0.003871	0.0	0.000	0.002	0.006	0.008
fetal_movement	2113.0	0.009517	0.046804	0.0	0.000	0.000	0.003	0.400
uterine_contractions	2113.0	0.004387	0.002941	0.0	0.002	0.005	0.007	0.008
light_decelerations	2113.0	0.001901	0.002966	0.0	0.000	0.000	0.003	0.008
severe_decelerations	2113.0	0.000003	0.000057	0.0	0.000	0.000	0.000	0.008
prolongued_decelerations	2113.0	0.000159	0.000592	0.0	0.000	0.000	0.000	0.008
abnormal_short_term_variability	2113.0	46.993848	17.177782	12.0	32.000	49.000	61.000	87.000
mean_value_of_short_term_variability	2113.0	1.335021	0.884368	0.2	0.700	1.200	1.700	7.000
percentage_of_time_with_abnormal_long_term_variability	2113.0	9.795078	18.337073	0.0	0.000	0.000	11.000	91.000
mean_value_of_long_term_variability	2113.0	8.166635	5.632912	0.0	4.600	7.400	10.800	50.700
histogram_width	2113.0	70.535258	39.007706	3.0	37.000	68.000	100.000	180.000
histogram_min	2113.0	93.564600	29.562269	50.0	67.000	93.000	120.000	159.000
histogram_max	2113.0	164.099858	17.945175	122.0	152.000	162.000	174.000	238.000
histogram_number_of_peaks	2113.0	4.077142	2.951664	0.0	2.000	4.000	6.000	18.000
histogram_number_of_zeroes	2113.0	0.325603	0.707771	0.0	0.000	0.000	0.000	10.000
histogram_mode	2113.0	137.454330	16.402026	60.0	129.000	139.000	148.000	187.000
histogram_mean	2113.0	134.599621	15.610422	73.0	125.000	136.000	145.000	182.000
histogram_median	2113.0	138.089446	14.478957	77.0	129.000	139.000	148.000	186.000
histogram_variance	2113.0	18.907241	29.038766	0.0	2.000	7.000	24.000	269.000
histogram_tendency	2113.0	0.318504	0.611075	-1.0	0.000	0.000	1.000	1.000
fetal_health	2113.0	1.303833	0.614279	1.0	1.000	1.000	1.000	3.000

Visualizing Fetal Health Classification Raw Data

In []:

```
normal = data.loc[data['fetal_health'] == 1].shape[0]
suspect = data.loc[data['fetal_health'] == 2].shape[0]
pathological = data.loc[data['fetal_health'] == 3].shape[0]
total = data.fetal_health.shape[0]

print("Total count: " + str(total))
print("Normal count: " + str(normal))
print("Suspect count: " + str(suspect))
print("Pathological count: " + str(pathological))


plt.figure(figsize = (10,5))
pie_fetal_health = plt.pie([normal, suspect, pathological], labels=["Normal", "Suspect", "Pathological"], autopct="%1.0f%%")
plt.title("Fetal health count")

plt.figure(figsize = (10,5))
data['fetal_health'].value_counts().plot(figsize=(10, 5), kind="bar")
plt.title("Fetal health count")
plt.xlabel("Fetal health score")
plt.ylabel("Frequency")

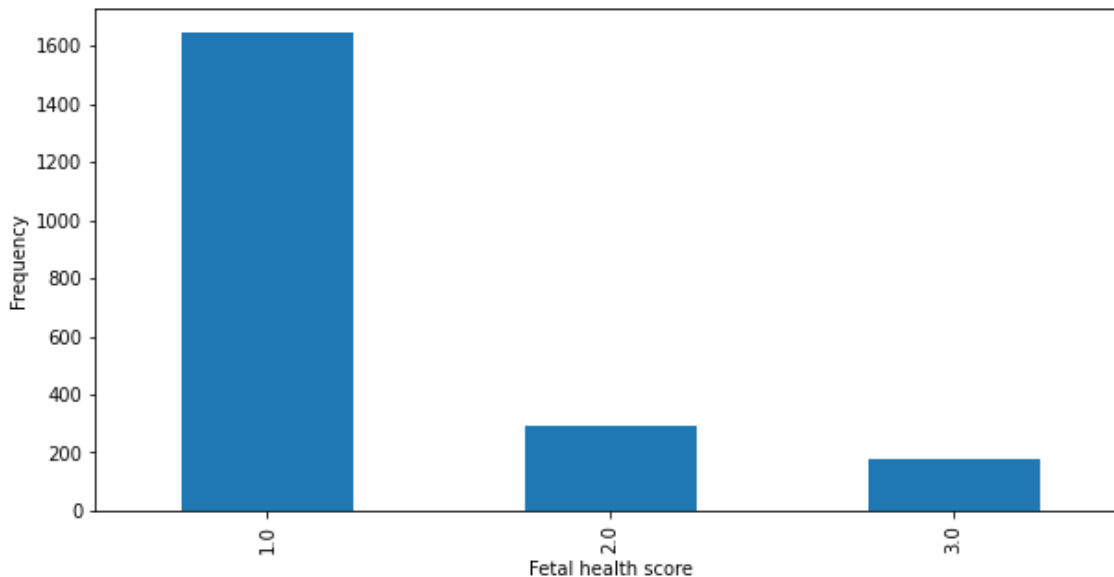
plt.show()
```

Total count: 2113
Normal count: 1646
Suspect count: 292
Pathological count: 175

Fetal health count



Fetal health count

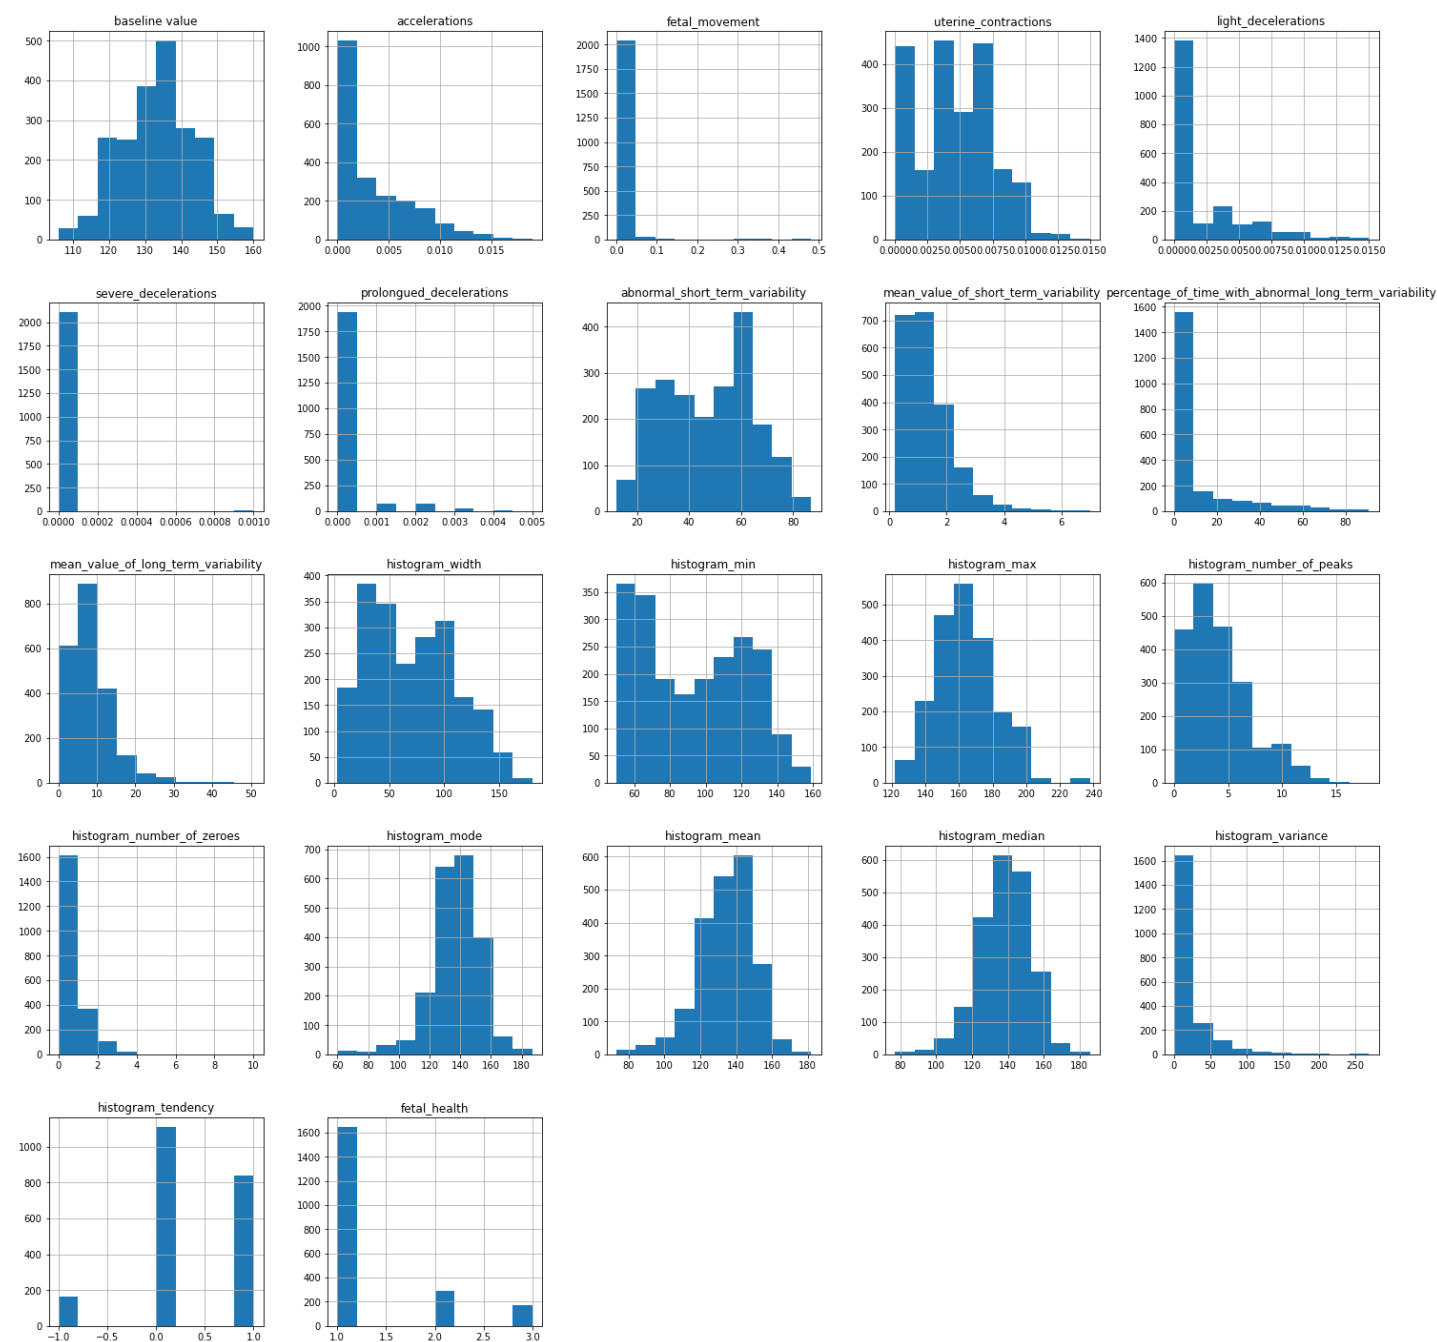


Note the imbalanced dataset in terms of fetal health scoring, leading to classification inaccuracy. To better observe feature importance and correlation, a confusion matrix will be assembled to observe correlation coefficients, giving us a better idea of what the most important features are.

Feature Analysis and Selection

In []:

```
feature_hist_plot = data.hist(figsize = (25,25))
```




Confusion Matrix

In []:

```
plt.figure(figsize=(20, 20))
correlation_matrix = sns.heatmap(data.corr(), annot=True, cmap="Purples")
```





```
In [ ]:
feature_correlation = data.corr()["fetal_health"].sort_values(ascending=False).to_frame()
feature_correlation.style.background_gradient(cmap=cm.Blues)
```

Out []:

	fetal_health
fetal_health	1.000000
prolongued_decelerations	0.486752
abnormal_short_term_variability	0.469671
percentage_of_time_with_abnormal_long_term_variability	0.421634
histogram_variance	0.208171
baseline_value	0.146077
severe_decelerations	0.132408
fetal_movement	0.088057
histogram_min	0.063529
light_decelerations	0.059651
histogram_number_of_zeroes	-0.016376
histogram_number_of_peaks	-0.022856
histogram_max	-0.046480
histogram_width	-0.069529

mean_value_of_short_term_variability	fetal_health
histogram_tendency	-0.135573
uterine_contractions	-0.205117
histogram_median	-0.208334
mean_value_of_long_term_variability	-0.225685
histogram_mean	-0.230243
histogram_mode	-0.253612
accelerations	-0.363947

Note that the row we are focused on is that of "fetal_health", where we can clearly see which features are most correlated with fetal_health. Of the 21 features, the features with highest correlation is "prolongued_decelerations, "abnormal_short_term_variability", and "percentage_of_time_with_abnormal_long_term_variability".

Now that we have our correlation coefficients, let's select the most important features with the KBest algorithm.

```
In [ ]:
X = data.drop(["fetal_health"], axis=1)
Y = data["fetal_health"]
selected_features = SelectKBest(score_func=f_classif, k='all')

feature_fit = selected_features.fit(X, Y)
feature_scores = pd.DataFrame({"Features": X.columns,
                              "Scores": feature_fit.scores_})


feature_scores
```

Out[]:

	Features	Scores
0	baseline value	137.833999
1	accelerations	194.618345
2	fetal_movement	11.700712
3	uterine_contractions	93.647474
4	light_decelerations	66.750344
5	severe_decelerations	28.438837
6	prolongued_decelerations	507.304309
7	abnormal_short_term_variability	337.703020
8	mean_value_of_short_term_variability	118.050463
9	percentage_of_time_with_abnormal_long_term_var...	335.386156
10	mean_value_of_long_term_variability	69.418940
11	histogram_width	54.215605
12	histogram_min	86.468440
13	histogram_max	2.523350
14	histogram_number_of_peaks	11.726828
15	histogram_number_of_zeroes	2.134901
16	histogram_mode	276.382795
17	histogram_mean	298.759569
18	histogram_median	249.699523
19	histogram_variance	150.955827
20	histogram_tendency	44.854186

In []:

```
plt.figure(figsize = (15,10))
plot = sns.barplot(data=feature_scores, x='Scores', y='Features', palette="rocket")
```



Let's now select the features with scores above the 100 threshold as our most important features and the ones that we will train our models on.

In []:

```
best_features = feature_scores[feature_scores['Scores'] > 100]
best_features
```

Out[]:

	Features	Scores
0	baseline value	137.833999
1	accelerations	194.618345
6	prolongued_decelerations	507.304309
7	abnormal_short_term_variability	337.703020
8	mean_value_of_short_term_variability	118.050463
9	percentage_of_time_with_abnormal_long_term_var...	335.386156
16	histogram_mode	276.382795
17	histogram_mean	298.759569
18	histogram_median	249.699523
19	histogram_variance	150.955827

In []:

```
best_features_array = list(best_features['Features'])
best_features_array.append('fetal_health')
data = data[best_features_array]
data.head()
```

Out[]:

baseline value accelerations prolonged_decelerations abnormal_short_term_variability mean_value_of_short_term_variability per

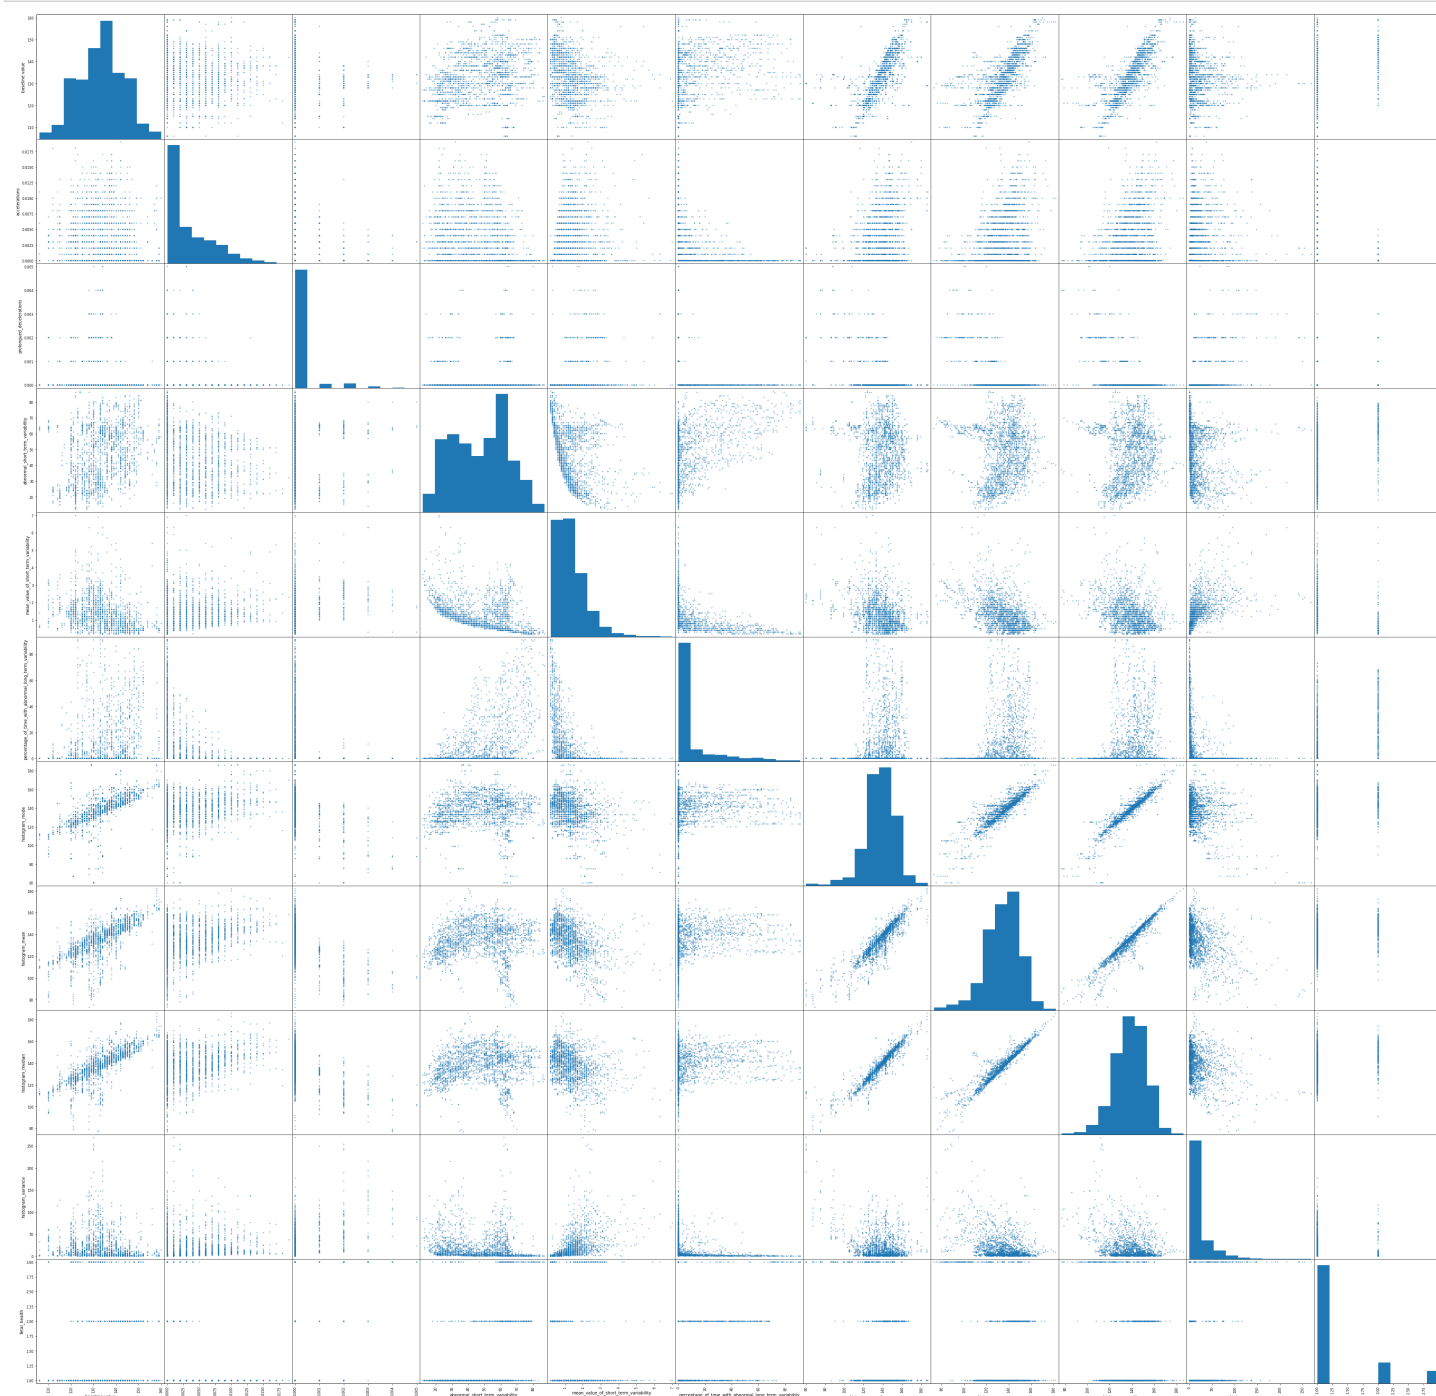
0	baseline_value	acceleration	prolonged_deceleration	abnormal_short_term_variability	mean_value_of_short_term_variability	per
1	132.0	0.006	0.0	17.0	2.1	
2	133.0	0.003	0.0	16.0	2.1	
3	134.0	0.003	0.0	16.0	2.4	
4	132.0	0.007	0.0	16.0	2.4	

Scatter Matrix

Visualizes the relationships between the most important features

In []:

```
from pandas.plotting import scatter_matrix
matrix = scatter_matrix(data, figsize=(60, 60))
```



Dataset split (train/test)

In order to use a model, we must first scale the array of features so that they are in the same range. We will also be using the standard 70/30 train/test split.

In []:

```
scaler = StandardScaler()
X=data.drop(["fetal_health"], axis=1)
X_df = pd.DataFrame(scaler.fit_transform(X), columns=X.columns)
X_df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 2113 entries, 0 to 2112
Data columns (total 10 columns):
#   Column                                     Non-Null Count  Dtype
---  -
0   baseline value                           2113 non-null   float64
1   accelerations                           2113 non-null   float64
2   prolonged_decelerations                 2113 non-null   float64
3   abnormal_short_term_variability         2113 non-null   float64
4   mean_value_of_short_term_variability    2113 non-null   float64
5   percentage_of_time_with_abnormal_long_term_variability 2113 non-null   float64
6   histogram_mode                         2113 non-null   float64
7   histogram_mean                         2113 non-null   float64
8   histogram_median                       2113 non-null   float64
9   histogram_variance                     2113 non-null   float64
dtypes: float64(10)
memory usage: 165.2 KB
```

In []:

```
X_df.head()
```

Out[]:

	baseline value	accelerations	prolongued_decelerations	abnormal_short_term_variability	mean_value_of_short_term_variability	percentage_of_time_with_abnormal_long_term_variability
0	1.352782	-0.823776	-0.26964	1.514300	-0.944425	0.865205
1	0.132665	0.726444	-0.26964	-1.746497	0.865205	0.865205
2	0.030989	-0.048666	-0.26964	-1.804726	0.865205	0.865205
3	0.070687	-0.048666	-0.26964	-1.804726	1.204511	0.865205
4	0.132665	0.984814	-0.26964	-1.804726	1.204511	0.865205

In []:

```
X_df.describe().T
```

Out[]:

	count	mean	std	min	25%	50%	75%
baseline value	2113.0	5.880031e-16	1.000237	2.776252	0.742724	0.030989	0.680746
accelerations	2113.0	9.678328e-17	1.000237	0.823776	0.823776	0.307036	0.726444
prolongued_decelerations	2113.0	2.528818e-15	1.000237	0.269640	0.269640	0.269640	0.269640
abnormal_short_term_variability	2113.0	6.567812e-17	1.000237	2.037640	0.873069	0.116815	0.815557
mean_value_of_short_term_variability	2113.0	3.913365e-16	1.000237	-	-	-	0.412708

	mean_value_of_short_term_variability	count	mean	std	min	25%	50%	75%
percentage_of_time_with_abnormal_long_term_variability	2113.0	2.212985e-15	1.000237	-	0.534294	0.534294	0.534294	0.065725
histogram_mode	2113.0	3.792255e-16	1.000237	-	4.723360	0.515566	0.094259	0.643101
histogram_mean	2113.0	7.585035e-16	1.000237	-	3.946992	0.615095	0.089729	0.666404
histogram_median	2113.0	9.018920e-16	1.000237	-	4.220187	0.627918	0.062903	0.684642
histogram_variance	2113.0	3.865026e-16	1.000237	-	0.651258	0.582368	0.410143	0.175419

In []:

```
y=data["fetal_health"]
X_train, X_test, y_train, y_test = train_test_split(X_df, y, test_size=0.3, random_state=42)
print("X_train samples: " + str(X_train.shape[0]) + " with " + str(X_train.shape[1]) + " features")
print("Y_train samples: " + str(y_train.shape[0]))
print("X_test samples: " + str(X_test.shape[0]) + " with " + str(X_test.shape[1]) + " features")
print("Y_test samples: " + str(y_test.shape[0]))
```

X_train samples: 1479 with 10 features
Y_train samples: 1479
X_test samples: 634 with 10 features
Y_test samples: 634

Machine Learning Classification

Here we will be using several classifiers to compare accuracy scores - we will be settling and fine-tuning the best classifier after comparing effectiveness.

Logistic Regression

In []:

```
from sklearn.linear_model import LogisticRegression
```

In []:

```
log_reg = LogisticRegression()
log_reg.fit(X_train, y_train)
log_reg_score = log_reg.score(X_test, y_test)
print("Baseline LR Score: " + str(log_reg_score))
```

Baseline LR Score: 0.8958990536277602

In []:

```
cv_score_lr = cross_val_score(log_reg, X_train, y_train)
print("Baseline LR CV Score: " + str(cv_score_lr.mean()))
```

Baseline LR CV Score: 0.8823751717819514

In []:

```
lr_predictions = log_reg.predict(X_test)
print(classification_report(y_test, lr_predictions))
```


	precision	recall	f1-score	support
1.0	0.93	0.97	0.95	497
2.0	0.73	0.58	0.65	90
3.0	0.78	0.77	0.77	47
accuracy			0.90	634
macro avg	0.81	0.77	0.79	634
weighted avg	0.89	0.90	0.89	634

In []:

```
plt.subplots(figsize=(12,8))
confusion_matrix_lr = confusion_matrix(y_test, lr_predictions)
sns.heatmap(confusion_matrix_lr/np.sum(confusion_matrix_lr),annot = True, cmap="Purples")
```

Out[]:

<matplotlib.axes._subplots.AxesSubplot at 0x7f4af10abad0>



Random Forest

In []:

```
from sklearn.ensemble import RandomForestClassifier
```

In []:

```
random_forest = RandomForestClassifier()
random_forest.fit(X_train, y_train)
random_forest_score = random_forest.score(X_test, y_test)
print("Baseline Random Forest Score: " + str(random_forest_score))
```

Baseline Random Forest Score: 0.9495268138801262

In []:

```
cv_score_rf = cross_val_score(random_forest, X_train, y_train)
print("Baseline RF CV Score: " + str(cv_score_rf.mean()))
```

Baseline RF CV Score: 0.9323797526339899

In []:

```
rf_predictions = random_forest.predict(X_test)
print(classification_report(y_test, rf_predictions))
```


	precision	recall	f1-score	support
1.0	0.96	0.99	0.97	497
2.0	0.91	0.74	0.82	90
3.0	0.88	0.96	0.92	47
accuracy			0.95	634
macro avg	0.92	0.90	0.90	634
weighted avg	0.95	0.95	0.95	634

In []:

```
plt.subplots(figsize=(12,8))
confusion_matrix_rf = confusion_matrix(y_test, rf_predictions)
sns.heatmap(confusion_matrix_rf/np.sum(confusion_matrix_rf),annot = True, cmap="Purples")
```

Out[]:

<matplotlib.axes._subplots.AxesSubplot at 0x7f4ae0b95f90>



K-Nearest Neighbors

In []:

```
from sklearn.neighbors import KNeighborsClassifier
```

In []:

```
knn = KNeighborsClassifier()
knn.fit(X_train, y_train)
knn_score = knn.score(X_test, y_test)
print("Baseline KNN Score: " + str(knn_score))
```

Baseline KNN Score: 0.9400630914826499

In []:

```
cv_score_knn = cross_val_score(knn, X_train, y_train)

print("Baseline KNN CV Score: " + str(cv_score_knn.mean()))
```

Baseline KNN CV Score: 0.9060215300045809

In []:

```
knn_predictions = knn.predict(X_test)
print(classification_report(y_test, knn_predictions))
```


	precision	recall	f1-score	support
1.0	0.95	0.98	0.97	497
2.0	0.88	0.71	0.79	90
3.0	0.91	0.91	0.91	47
accuracy			0.94	634
macro avg	0.91	0.87	0.89	634
weighted avg	0.94	0.94	0.94	634

In []:

```
plt.subplots(figsize=(12,8))
confusion_matrix_knn = confusion_matrix(y_test, knn_predictions)
sns.heatmap(confusion_matrix_knn/np.sum(confusion_matrix_knn),annot = True, cmap="Purples")
```

Out[]:

<matplotlib.axes._subplots.AxesSubplot at 0x7f4ae0983690>



Support Vector Classifier

In []:

```
from sklearn.svm import SVC
```

In []:

```
svc = SVC()
svc.fit(X_train, y_train)
svc_score = svc.score(X_test, y_test)
print("Baseline SVC Score: " + str(svc_score))
```

Baseline SVC Score: 0.9211356466876972

In []:

```
vc_score_svc = cross_val_score(svc, X_train, y_train)
print("Baseline SVC CV Score: " + str(vc_score_svc.mean()))
```

Baseline SVC CV Score: 0.8924965643609711

In []:

```
svc_predictions = svc.predict(X_test)
print(classification_report(y_test, svc_predictions))
```


	precision	recall	f1-score	support
1.0	0.94	0.97	0.95	497
2.0	0.78	0.67	0.72	90
3.0	0.98	0.89	0.93	47
accuracy			0.92	634
macro avg	0.90	0.84	0.87	634
weighted avg	0.92	0.92	0.92	634

In []:

```
plt.subplots(figsize=(12,8))
confusion_matrix_svc = confusion_matrix(y_test, svc_predictions)
sns.heatmap(confusion_matrix_svc/np.sum(confusion_matrix_svc),annot = True, cmap="Purples")
```

Out[]:

<matplotlib.axes._subplots.AxesSubplot at 0x7f4ae08af050>



Linear Support Vector Classifier

In []:

```
from sklearn.svm import LinearSVC
```

In []:

```
linear_svc = LinearSVC()
linear_svc.fit(X_train, y_train)
linear_svc_score = linear_svc.score(X_test, y_test)
print("Baseline Linear SVC Score: " + str(linear_svc_score))
```

Baseline Linear SVC Score: 0.8927444794952681

/usr/local/lib/python3.7/dist-packages/sklearn/svm/_base.py:947: ConvergenceWarning: Liblinear failed to converge, increase the number of iterations.
"the number of iterations.", ConvergenceWarning)

In []:

```
cv_score_linear_svc = cross_val_score(linear_svc, X_train, y_train)

print("Baseline Linear SVC CV Score: " + str(cv_score_linear_svc.mean()))
```

/usr/local/lib/python3.7/dist-packages/sklearn/svm/_base.py:947: ConvergenceWarning: Liblinear failed to converge, increase the number of iterations.
"the number of iterations.", ConvergenceWarning)
/usr/local/lib/python3.7/dist-packages/sklearn/svm/_base.py:947: ConvergenceWarning: Liblinear failed to converge, increase the number of iterations.
"the number of iterations.", ConvergenceWarning)
/usr/local/lib/python3.7/dist-packages/sklearn/svm/_base.py:947: ConvergenceWarning: Liblinear failed to converge, increase the number of iterations.
"the number of iterations.", ConvergenceWarning)

Baseline Linear SVC CV Score: 0.8843999083829592

/usr/local/lib/python3.7/dist-packages/sklearn/svm/_base.py:947: ConvergenceWarning: Liblinear failed to converge, increase the number of iterations.
"the number of iterations.", ConvergenceWarning)
/usr/local/lib/python3.7/dist-packages/sklearn/svm/_base.py:947: ConvergenceWarning: Liblinear failed to converge, increase the number of iterations.
"the number of iterations.", ConvergenceWarning)

In []:

```
linear_svc_predictions = linear_svc.predict(X_test)
print(classification_report(y_test, linear_svc_predictions))
```


	precision	recall	f1-score	support
1.0	0.92	0.97	0.94	497
2.0	0.73	0.53	0.62	90
3.0	0.86	0.79	0.82	47
accuracy			0.89	634
macro avg	0.83	0.76	0.79	634
weighted avg	0.89	0.89	0.89	634

In []:

```
plt.subplots(figsize=(12,8))
confusion_matrix_linear_svc = confusion_matrix(y_test, linear_svc_predictions)
sns.heatmap(confusion_matrix_linear_svc/np.sum(confusion_matrix_linear_svc),annot = True
, cmap="Purples")
```

Out []:

<matplotlib.axes._subplots.AxesSubplot at 0x7f4ae07ec2d0>



Decision Tree

In []:

```
from sklearn.tree import DecisionTreeClassifier
```

In []:

```
decision_tree = DecisionTreeClassifier()
decision_tree.fit(X_train, y_train)
decision_tree_score = decision_tree.score(X_test, y_test)
print("Baseline Decision Tree Score: " + str(decision_tree_score))
```

Baseline Decision Tree Score: 0.9337539432176656

In []:

```
cv_score_decision_tree = cross_val_score(decision_tree, X_train, y_train)
print("Baseline Decision Tree CV Score: " + str(cv_score_decision_tree.mean()))
```

Baseline Decision Tree CV Score: 0.9073522675217591

In []:

```
decision_tree_predictions = decision_tree.predict(X_test)
print(classification_report(y_test, decision_tree_predictions))
```


	precision	recall	f1-score	support
1.0	0.96	0.97	0.96	497
2.0	0.80	0.73	0.77	90
3.0	0.88	0.96	0.92	47
accuracy			0.93	634
macro avg	0.88	0.89	0.88	634
weighted avg	0.93	0.93	0.93	634

In []:


```
plt.subplots(figsize=(12,8))
confusion_matrix_decision_tree = confusion_matrix(y_test, decision_tree_predictions)
sns.heatmap(confusion_matrix_decision_tree/np.sum(confusion_matrix_decision_tree),annot =
True, cmap="Purples")
```

Out[]:

<matplotlib.axes._subplots.AxesSubplot at 0x7f4ae0719b90>



Multi-layer Perceptron (Neural Network)

In []:

```
from sklearn.neural_network import MLPClassifier
```

In []:

```
neural_network = MLPClassifier()
neural_network.fit(X_train, y_train)
neural_network_score = neural_network.score(X_test, y_test)
print("Baseline Multi-layer Perceptron Score: " + str(neural_network_score))
```

Baseline Multi-layer Perceptron Score: 0.9353312302839116

```
/usr/local/lib/python3.7/dist-packages/sklearn/neural_network/_multilayer_perceptron.py:5
71: ConvergenceWarning: Stochastic Optimizer: Maximum iterations (200) reached and the op
timization hasn't converged yet.
% self.max_iter, ConvergenceWarning)
```

In []:

```
cv_score_neural_network = cross_val_score(neural_network, X_train, y_train)
print("Baseline Neural Network CV Score: " + str(cv_score_neural_network.mean()))
```

```
/usr/local/lib/python3.7/dist-packages/sklearn/neural_network/_multilayer_perceptron.py:5
71: ConvergenceWarning: Stochastic Optimizer: Maximum iterations (200) reached and the op
timization hasn't converged yet.
% self.max_iter, ConvergenceWarning)
/usr/local/lib/python3.7/dist-packages/sklearn/neural_network/_multilayer_perceptron.py:5
71: ConvergenceWarning: Stochastic Optimizer: Maximum iterations (200) reached and the op
timization hasn't converged yet.
```

```

% self.max_iter, ConvergenceWarning)
/usr/local/lib/python3.7/dist-packages/sklearn/neural_network/_multilayer_perceptron.py:571: ConvergenceWarning: Stochastic Optimizer: Maximum iterations (200) reached and the optimization hasn't converged yet.
% self.max_iter, ConvergenceWarning)
/usr/local/lib/python3.7/dist-packages/sklearn/neural_network/_multilayer_perceptron.py:571: ConvergenceWarning: Stochastic Optimizer: Maximum iterations (200) reached and the optimization hasn't converged yet.
% self.max_iter, ConvergenceWarning)

```

Baseline Neural Network CV Score: 0.9107627118644068

```

/usr/local/lib/python3.7/dist-packages/sklearn/neural_network/_multilayer_perceptron.py:571: ConvergenceWarning: Stochastic Optimizer: Maximum iterations (200) reached and the optimization hasn't converged yet.
% self.max_iter, ConvergenceWarning)

```

In []:

```

neural_network_predictions = neural_network.predict(X_test)
print(classification_report(y_test, neural_network_predictions))

```

	precision	recall	f1-score	support
1.0	0.96	0.97	0.96	497
2.0	0.81	0.76	0.78	90
3.0	0.90	0.94	0.92	47
accuracy			0.94	634
macro avg	0.89	0.89	0.89	634
weighted avg	0.93	0.94	0.93	634

In []:


```

plt.subplots(figsize=(12,8))
neural_network_decision_tree = confusion_matrix(y_test, neural_network_predictions)
sns.heatmap(neural_network_decision_tree/np.sum(neural_network_decision_tree),annot = True, cmap="Purples")

```

Out[]:

<matplotlib.axes._subplots.AxesSubplot at 0x7f4ae065c310>



Gradient Boosting

In []:

```
from sklearn.ensemble import GradientBoostingClassifier
```

In []:

```
gradient_boosting = GradientBoostingClassifier()
gradient_boosting.fit(X_train, y_train)
gradient_boosting_score = gradient_boosting.score(X_test, y_test)
print("Baseline Gradient Boosting Classifier Score: " + str(gradient_boosting_score))
```

Baseline Gradient Boosting Classifier Score: 0.9558359621451105

In []:

```
cv_score_gb = cross_val_score(gradient_boosting, X_train, y_train)
print("Baseline GB CV Score: " + str(cv_score_gb.mean()))
```

Baseline GB CV Score: 0.9371117727897389

In []:

```
gb_predictions = gradient_boosting.predict(X_test)
print(classification_report(y_test, gb_predictions))
```


	precision	recall	f1-score	support
1.0	0.97	0.98	0.98	497
2.0	0.90	0.80	0.85	90
3.0	0.92	0.96	0.94	47
accuracy			0.96	634
macro avg	0.93	0.91	0.92	634
weighted avg	0.95	0.96	0.95	634


In []:

```
plt.subplots(figsize=(12,8))
gb_decision_tree = confusion_matrix(y_test, gb_predictions)
sns.heatmap(gb_decision_tree/np.sum(gb_decision_tree),annot = True, cmap="Purples")
```

Out[]:

<matplotlib.axes._subplots.AxesSubplot at 0x7f4ae05e90d0>





LightGBM Classifier

In []:

```
from lightgbm import LGBMClassifier
```

In []:

```
light_gbm = LGBMClassifier()
light_gbm.fit(X_train, y_train)
light_gbm_score = light_gbm.score(X_test, y_test)
print("Baseline LGBM Classifier Score: " + str(light_gbm_score))
```

Baseline LGBM Classifier Score: 0.9558359621451105

In []:

```
cv_score_light_gbm = cross_val_score(light_gbm, X_train, y_train)

print("Baseline LightGBM CV Score: " + str(cv_score_light_gbm.mean()))
```

Baseline LightGBM CV Score: 0.9384722858451671

In []:

```
light_gbm_predictions = light_gbm.predict(X_test)
print(classification_report(y_test, light_gbm_predictions))
```

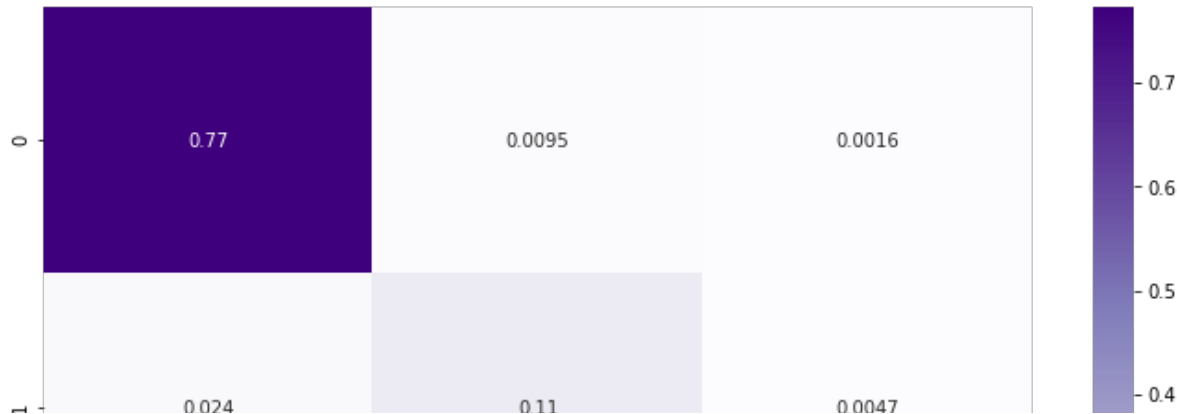
	precision	recall	f1-score	support
1.0	0.97	0.99	0.98	497
2.0	0.90	0.80	0.85	90
3.0	0.92	0.94	0.93	47
accuracy			0.96	634
macro avg	0.93	0.91	0.92	634
weighted avg	0.95	0.96	0.95	634


In []:

```
plt.subplots(figsize=(12,8))
light_gbm_decision_tree = confusion_matrix(y_test, light_gbm_predictions)
sns.heatmap(light_gbm_decision_tree/np.sum(light_gbm_decision_tree),annot = True, cmap="Purples")
```

Out[]:

<matplotlib.axes._subplots.AxesSubplot at 0x7f4ae053fad0>





XGBoost Classifier

In []:

```
from xgboost import XGBClassifier
import xgboost
```

In []:

```
xgboost = XGBClassifier()
xgboost.fit(X_train, y_train)
xgboost_score = xgboost.score(X_test, y_test)
print("Baseline XGBoost Classifier Score: " + str(xgboost_score))
```

Baseline XGBoost Classifier Score: 0.9495268138801262

In []:

```
cv_score_xgboost = cross_val_score(xgboost, X_train, y_train)
print("Baseline XGBoost CV Score: " + str(cv_score_xgboost.mean()))
```

Baseline XGBoost CV Score: 0.9357581310123683

In []:

```
xgboost_predictions = xgboost.predict(X_test)
print(classification_report(y_test, xgboost_predictions))
```


	precision	recall	f1-score	support
1.0	0.97	0.98	0.97	497
2.0	0.87	0.80	0.83	90
3.0	0.92	0.94	0.93	47
accuracy			0.95	634
macro avg	0.92	0.90	0.91	634
weighted avg	0.95	0.95	0.95	634

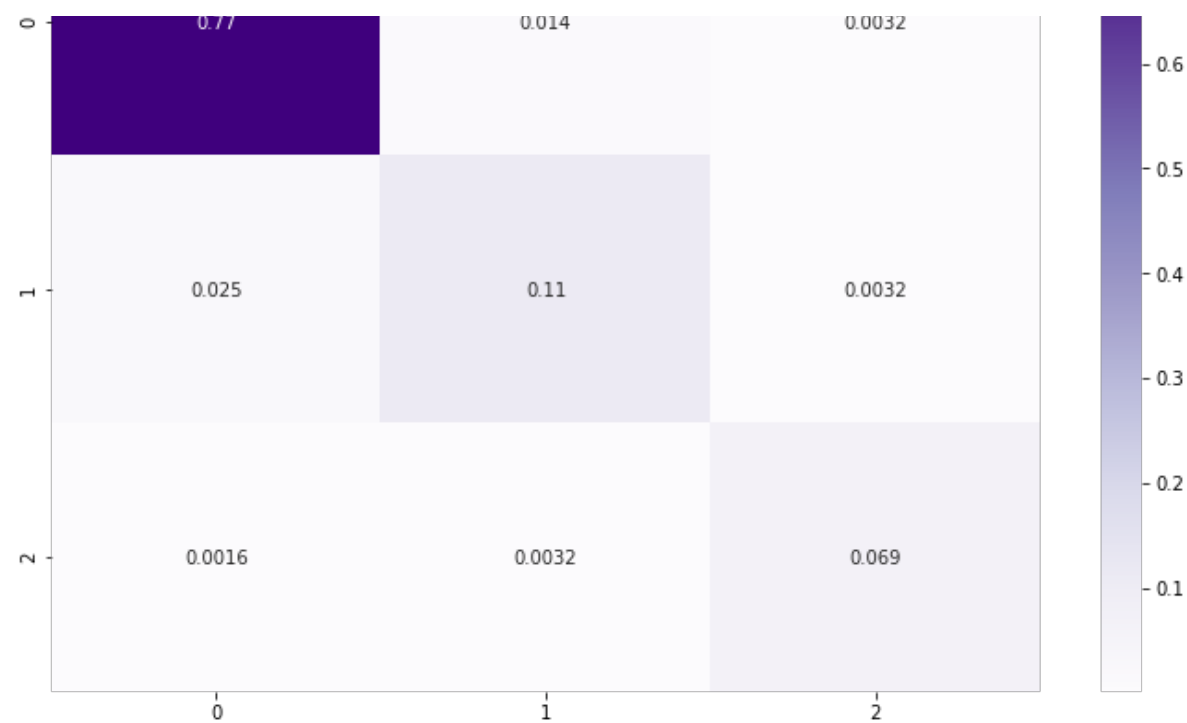
In []:

```
plt.subplots(figsize=(12,8))
xgboost_decision_tree = confusion_matrix(y_test, xgboost_predictions)
sns.heatmap(xgboost_decision_tree/np.sum(xgboost_decision_tree),annot = True, cmap="Purples")
```

Out []:

<matplotlib.axes._subplots.AxesSubplot at 0x7f4ae0478550>





Model Scoring and Selection

```
In [ ]:

model_scores = pd.DataFrame({"Model": ["Logistic Regression", "Random Forest",
                                         "K-Nearest Neighbors", "Support Vector", "Linear S
                                         "Decision Tree", "Multi-layer Perceptron", "Gradie
                                         "LightGBM Classifier", "XGBoost Classifier"],
                              "Scores": [log_reg_score, random_forest_score, knn_score, sv
                                         c_score, linear_svc_score, decision_tree_score,
                                         neural_network_score, gradient_boosting_score, li
                                         ght_gbm_score, xgboost_score]})
model_scores.sort_values(by="Scores", ascending=False)
```

Out[]:

	Model	Scores
7	Gradient Boosting	0.955836
8	LightGBM Classifier	0.955836
1	Random Forest	0.949527
9	XGBoost Classifier	0.949527
2	K-Nearest Neighbors	0.940063
6	Multi-layer Perceptron	0.935331
5	Decision Tree	0.933754
3	Support Vector	0.921136
0	Logistic Regression	0.895899
4	Linear Support Vector	0.892744

Model Tuning and Optimization

LightGBM Classifier Tuning/Optimization

```
In [ ]:
```

```
param_grid = {
    'learning_rate': [0.001,0.01],
    'n_estimators': [ 1000],
    'num_leaves': [12, 30,80],
    'boosting_type' : ['gbdt'],
    'objective' : ['binary'],
    'random_state' : [1],
    'colsample_bytree' : [ 0.8, 1],
    'subsample' : [0.5,0.7,0.75],
    'reg_alpha' : [0.1, 1.2],
    'reg_lambda' : [0.1, 1.2],
    'subsample_freq' : [500,1000],
    'max_depth' : [15, 30, 80]
}
```

In []:

```
cv_method = StratifiedKFold(n_splits=3)
```

In []:

```
GridSearchCV_GBM = GridSearchCV(estimator=LGBMClassifier(),
                                param_grid=param_grid,
                                cv=cv_method,
                                verbose=1,
                                n_jobs=4,
                                scoring="accuracy",
                                return_train_score=True
                                )
```

In []:

```
GridSearchCV_GBM.fit(X_train, y_train)
```

Fitting 3 folds for each of 864 candidates, totalling 2592 fits

```
[Parallel(n_jobs=4)]: Using backend LokyBackend with 4 concurrent workers.
[Parallel(n_jobs=4)]: Done 42 tasks      | elapsed: 36.7s
[Parallel(n_jobs=4)]: Done 192 tasks    | elapsed: 2.8min
[Parallel(n_jobs=4)]: Done 442 tasks    | elapsed: 6.5min
[Parallel(n_jobs=4)]: Done 792 tasks    | elapsed: 11.8min
[Parallel(n_jobs=4)]: Done 1242 tasks   | elapsed: 19.4min
[Parallel(n_jobs=4)]: Done 1792 tasks   | elapsed: 27.9min
[Parallel(n_jobs=4)]: Done 2442 tasks   | elapsed: 38.9min
[Parallel(n_jobs=4)]: Done 2592 out of 2592 | elapsed: 41.6min finished
```

Out[]:

```
GridSearchCV(cv=StratifiedKFold(n_splits=3, random_state=None, shuffle=False),
            error_score=nan,
            estimator=LGBMClassifier(boosting_type='gbdt', class_weight=None,
                                     colsample_bytree=1.0,
                                     importance_type='split',
                                     learning_rate=0.1, max_depth=-1,
                                     min_child_samples=20,
                                     min_child_weight=0.001,
                                     min_split_gain=0.0, n_estimators=100,
                                     n_jobs=-1, num_leaves=31, objective=None...
            'colsample_bytree': [0.8, 1],
            'learning_rate': [0.001, 0.01],
            'max_depth': [15, 30, 80], 'n_estimators': [1000],
            'num_leaves': [12, 30, 80], 'objective': ['binary'],
            'random_state': [1], 'reg_alpha': [0.1, 1.2],
            'reg_lambda': [0.1, 1.2],
            'subsample': [0.5, 0.7, 0.75],
            'subsample_freq': [500, 1000]},
            pre_dispatch='2*n_jobs', refit=True, return_train_score=True,
            scoring='accuracy', verbose=1)
```

In []:

```
best_params_GBM = GridSearchCV_GBM.best_params_
```

```
best_params_GBM = GridSearchCV_GBM.best_params_  
best_params_GBM
```

Out[]:

```
{'boosting_type': 'gbdt',  
 'colsample_bytree': 0.8,  
 'learning_rate': 0.01,  
 'max_depth': 15,  
 'n_estimators': 1000,  
 'num_leaves': 12,  
 'objective': 'binary',  
 'random_state': 1,  
 'reg_alpha': 0.1,  
 'reg_lambda': 1.2,  
 'subsample': 0.7,  
 'subsample_freq': 500}
```

In []:

```
best_score_GBM = GridSearchCV_GBM.best_score_  
print("Optimized Grid Search GBM best score: " + str(best_score_GBM))
```

Optimized Grid Search GBM best score: 0.9263015551048005

In []:

```
optimized_LGBM = LGBMClassifier(**GridSearchCV_GBM.best_params_)  
optimized_LGBM.fit(X_train, y_train)
```

Out[]:

```
LGBMClassifier(boosting_type='gbdt', class_weight=None, colsample_bytree=0.8,  
               importance_type='split', learning_rate=0.01, max_depth=15,  
               min_child_samples=20, min_child_weight=0.001, min_split_gain=0.0,  
               n_estimators=1000, n_jobs=-1, num_leaves=12, objective='binary',  
               random_state=1, reg_alpha=0.1, reg_lambda=1.2, silent=True,  
               subsample=0.7, subsample_for_bin=200000, subsample_freq=500)
```

In []:

```
optimized_LGBM_score = optimized_LGBM.score(X_test, y_test)  
print("Optimized LGBM Classifier Score: " + str(optimized_LGBM_score))
```

Optimized LGBM Classifier Score: 0.9479495268138801

In []:

```
cv_score_optimized_LGBM = cross_val_score(optimized_LGBM, X_train, y_train)  
print("Optimized LightGBM CV Score: " + str(cv_score_optimized_LGBM.mean()))
```

Optimized LightGBM CV Score: 0.9222446174988548

In []:

```
optimized_LGBM_predictions = optimized_LGBM.predict(X_test)  
print(classification_report(y_test, optimized_LGBM_predictions))
```

	precision	recall	f1-score	support
1.0	0.96	0.98	0.97	497
2.0	0.89	0.76	0.82	90
3.0	0.88	0.98	0.93	47
accuracy			0.95	634
macro avg	0.91	0.90	0.91	634
weighted avg	0.95	0.95	0.95	634

In []:

```
plt.subplots(figsize=(12,8))
```



```
optimized_light_gbm_decision_tree = confusion_matrix(y_test, optimized_LGBM_predictions)
sns.heatmap(optimized_light_gbm_decision_tree/np.sum(optimized_light_gbm_decision_tree),
            annot = True, cmap="Purples")
```

Out[]:

<matplotlib.axes._subplots.AxesSubplot at 0x7f4adf2c1a50>

